

1 MATRIX MULTIPLICATION

Given $A = (a_{ij})$ and $B = (b_{ij})$ are square matrices of order n , then $C = (c_{ij}) = AB$ is also a square matrix of order n , and c_{ij} is obtained by taking the dot product of the i th row of A with the j th column of B . In other words, $c_{ij} = a_{i0}b_{0j} + a_{i1}b_{1j} + \dots + a_{i,n-1}b_{n-1,j}$.

Matrix Multiplication Algorithm:

1. for each row of C
2. for each column of C
3. $C[\text{row}][\text{column}] = 0.0$;
4. for each element of this row of A
5. Add $A[\text{row}][\text{element}] * B[\text{element}][\text{column}]$ to $C[\text{row}][\text{column}]$

We can see that a straightforward parallel implementation of matrix multiplication is costly. If we suppose that the number of processes is the same as the order of the matrices; i.e, $p = n$, and suppose that we have distributed the matrices by rows. So process 0 is assigned row 0 of A , B , and C ; process 1 is assigned row 1 of A , B , and C ; etc. Then in order to form the dot product of the i th row of A with j th column of B , we will need to form the dot product of the j th column with every row of A . So we will have to carry out an allgather (gathers data from all tasks and distribute the combined data to all tasks) rather than a gather (gathers together values from a group of processes), and we will have to do this for every column of B . This involves a lot of expensive communication. Similarly, an algorithm that distributes the matrices by columns will involve large amount of communication. In order to reduce communication, most parallel matrix multiplication algorithms use a checkerboard distribution of the matrices. The matrices are then seen as a grid. Instead of assigning entire rows or columns to each process, we assign small submatrices. For example, if we have a 4×4 matrix, we can assign the elements of the matrix in the following way to 4 processes:

Process 0	Process 1
a_{00} a_{01}	a_{02} a_{03}
a_{10} a_{11}	a_{12} a_{13}
Process 2	Process 3
a_{20} a_{21}	a_{22} a_{23}
a_{30} a_{31}	a_{32} a_{33}

1.1 Fox's Algorithm

Assume that the matrices have order n , and the number of processes, p , equals n^2 . Then a checkerboard mapping assigns a_{ij} , b_{ij} , and c_{ij} to process (i, j) .

Fox's algorithm proceeds in n stages: one stage for each term $a_{ik}b_{kj}$ in the dot product

$$c_{ij} = a_{i0}b_{0j} + a_{i1}b_{1j} + \dots + a_{i,n-1}b_{n-1,j}.$$

1. Initial stage: each process multiplies the diagonal entry of A in its process row by its element of B :
Stage 0 on process (i, j) : $c_{ij} = a_{ii}b_{ij}$.
2. Next stage: each process multiplies the element immediately to the right of the diagonal of A (in its process row) by the process of B directly beneath its own element of B :
Stage 1 on process (i, j) : $c_{ij} = c_{ij} + a_{i,i+1}b_{i+1,j}$.
3. In general, during the k th stage, each process multiplies the element k columns to the right of the diagonal of A by the element k rows below its own element of B :
Stage k on process (i, j) : $c_{ij} = c_{ij} + a_{i,i+k}b_{i+k,j}$.
4. We may encounter out of range subscript when adding k to a row or column. Instead, we define $\bar{k} = (i+k) \bmod n$. Hence, for stage k on process (i, j) : $c_{ij} = c_{ij} + a_{i,\bar{k}}b_{\bar{k},j}$. c_{ij} will be computed as follows:
$$c_{ij} = a_{ii}b_{ij} + a_{i,i+1}b_{i+1,j} + \dots + a_{i,n-1}b_{n-1,j} + a_{i0}b_{0j} + \dots + a_{i,i-1}b_{i-1,j}.$$

We will store sub-matrices rather than matrix elements on each process, provided that the sub-matrices can be multiplied as needed and the number of process rows or process columns, \sqrt{p} , evenly divides n . Given this assumption, each process is assigned a square $\frac{n}{\sqrt{p}} \times \frac{n}{\sqrt{p}}$ sub-matrix of each of the three matrices. Define $q = \sqrt{p}$ and $\bar{n} = \frac{n}{\sqrt{p}}$. Denote the sub-matrices A_{ij} , B_{ij} , and C_{ij} by $A[i, j]$, $B[i, j]$, and $C[i, j]$, the **pseudo-code of Fox's algorithm is the following**:

1. my process row = i , my process column = j
2. $q = \sqrt{p}$
3. dest = $((i - 1) \bmod q, j)$
4. source = $((i + 1) \bmod q, j)$
5. for (stage = 0; stage < q ; stage++)
 6. $\bar{k} = (i + \text{stage}) \bmod q$
 7. Broadcast $A[i, \bar{k}]$ across process row i
 8. $C[i, j] = C[i, j] + A[i, \bar{k}] * B[\bar{k}, j]$
 9. Send $B[\bar{k}, j]$ to dest; Receive
 10. $B[(\bar{k} + 1) \bmod q, j]$ from source

The algorithm can also be written in the following way:

1. for (stage = 0; stage < q; stage++)
 2. Choose a sub-matrix of A of each row of processes
 3. In each row of processes broadcast the sub-matrix chosen in that row to the other processes in that row
 4. On each process, multiply the newly received sub-matrix of A by the sub-matrix of B currently residing on the process
 5. On each process, send the sub-matrix of B to the process directly above. (On processes in the first row, send the sub-matrix to the last row)

1.2 Example Applied to 2×2 Matrices

Assume we have n^2 processes, one for each element in A , B , and C .

Stage 0:

1. $A_{i,i}$ on process $P_{i,j}$, so broadcast the diagonal elements of A across the rows ($A_{ii} \rightarrow P_{ij}$). This will place $A_{0,0}$ on each $P_{0,j}$ and $A_{1,1}$ on each $P_{1,j}$.
The A elements on the P matrix will be:

A_{00}	A_{00}
A_{11}	A_{11}

2. $B_{i,j}$ on process $P_{i,j}$, so broadcast B across the rows ($B_{ij} \rightarrow P_{ij}$). The A and B values on the P matrix will be:

A_{00}	A_{00}
B_{00}	B_{01}
A_{11}	A_{11}
B_{10}	B_{11}

3. Compute $C_{ij} = AB$ for each processor:

A_{00}	A_{00}
B_{00}	B_{01}
$C_{00} = A_{00}B_{00}$	$C_{01} = A_{00}B_{01}$
A_{11}	A_{11}
B_{10}	B_{11}
$C_{10} = A_{11}B_{10}$	$C_{11} = A_{11}B_{11}$

Now we are ready for the next stage. In this stage, we broadcast the next column (mod n) of A across the processes and shift up (mod n) the B value.

Stage 1:

1. The next column is $A_{0,1}$ for the 1st row and $A_{1,0}$ for the 2nd row. Broadcast the next A across the rows:

A_{01}	A_{01}
B_{00}	B_{01}
$C_{00} = A_{00}B_{00}$	$C_{01} = A_{00}B_{01}$
A_{10}	A_{10}
B_{10}	B_{11}
$C_{10} = A_{11}B_{10}$	$C_{11} = A_{11}B_{11}$

2. Shift the B values up. $B_{1,0}$ moves up from process $P_{1,0}$ to process $P_{0,0}$ and $B_{0,0}$ moves up (mod n) from $P_{0,0}$ to $P_{1,0}$. Similarly for $B_{1,1}$ and $B_{0,1}$.

A_{01}	A_{01}
B_{10}	B_{11}
$C_{00} = A_{00}B_{00}$	$C_{01} = A_{00}B_{01}$
A_{10}	A_{10}
B_{00}	B_{01}
$C_{10} = A_{11}B_{10}$	$C_{11} = A_{11}B_{11}$

3. Compute $C_{ij} = AB$ for each process:

A_{01}	A_{01}
B_{10}	B_{11}
$C_{00} = C_{00} + A_{01}B_{01}$	$C_{01} = C_{01} + A_{01}B_{11}$
A_{10}	A_{10}
B_{00}	B_{01}
$C_{10} = C_{10} + A_{10}B_{00}$	$C_{11} = C_{11} + A_{10}B_{01}$

1.3 Cannon's Algorithm

In the following algorithm, all the subscripts are modulo n .

- Initially each $p_{i,j}$ has $a_{i,j}$ and $b_{i,j}$
- Align elements $a_{i,j}$ and $b_{i,j}$ by reordering so that $a_{i,j+1}$ and $b_{i+j,j}$ are on $p_{i,j}$
- Each $p_{i,j}$ computes $c_{i,j} = a_{i,j+1} \times b_{i+j,j}$ ($a_{i,j+1}$ and $b_{i+j,j}$ are local on $p_{i,j}$)
- For $k = 1$ to $n - 1$:
 - Rotate A left by one column
 - Rotate B up by one row
 - Each $p_{i,j}$ computes $c_{i,j} = c_{i,j} + a_{i,j+i+k} \times b_{i+j+k,j}$ ($a_{i,j+i+k}$ and $b_{i+j+k,j}$ are local on $p_{i,j}$ after k iterations).